

Why Does DeepSeek Pursue Alpha in Finance?

A research analyst's perspective on where AI and finance intersect

As of 2026, generative AI is used pervasively in investment research. So in this already-crowded market, why does DeepSeek emphasize financial reasoning, and can that translate into genuine excess returns (alpha)? This piece examines the argument not through benchmarks or marketing, but through the model's design architecture and the background of its parent company.

While running the [vibe-investing](#) repository and analyzing market conditions and individual stocks through various LLM prompts, I noticed something interesting: DeepSeek pursues the alpha returns of aggressive investing.

1. The Differentiator — Context, Not Calculation

The real weakness of financial AI is closer to "calculation without context" than to hallucination. Pulling figures like a forward P/E or the VIX can be done with a simple search. What matters is the ability to connect those numbers and convert them into the judgment of "is now the time to buy?"

DeepSeek combines an **MoE (Mixture of Experts)** architecture with **reinforcement-learning-based reasoning (Chain-of-Thought)**, designed so that a single model handles numerical computation, historical pattern matching, and counter-argumentation within one reasoning chain.

For example, in checking a correction-market buy signal, it moves through these stages: statistical contextualization ("a 36% gain is a hindsight figure; on a year-end buying basis, 19% is correct"), exception handling for historical cases (the failures of 1939, 1966, 1970, and 1974), and causal attribution (are interest rates, oil prices, and tariffs "political noise" or a "structural crisis"?).

In short, it is designed to mimic an analyst's thought process: hypothesis formation → cross-validation of data → attempt at refutation → confidence assignment.

2. The Structural Reason for Its Financial Strength: Parent Company High-Flyer

The core reason DeepSeek started from a different line than general-purpose LLMs lies in its corporate background. This model was not born in an academic lab or a Big Tech research division, but spun out of one of China's largest quantitative hedge funds, Huanfang Quant (幻方量化, High-Flyer Quant).

The nature of the parent company. Founder Liang Wenfeng (梁文锋), after graduating from Zhejiang University, co-founded the quantitative hedge fund High-Flyer in 2015 and introduced AI into its trading strategies. The entity was established in February 2016, and as of December 2025 it managed roughly USD 10 billion in assets. In 2025, High-Flyer managed about RMB 70 billion (roughly USD 10 billion) and posted an average return of 56.6%, ranking second among Chinese quant funds above RMB 10 billion in size. First place went to Lingjun Investment, which recorded 73.5%.

The conversion of compute infrastructure. Liang began buying thousands of NVIDIA GPUs from 2021, before the U.S. export restrictions on AI chips to China. These were initially for algorithmic trading, and later became the foundation for the 2023 launch of DeepSeek. The large-scale compute capacity and data-engineering know-how accumulated in financial operations were carried directly over into AI infrastructure.

The principle of independence. Until April 2026, DeepSeek operated entirely on High-Flyer's capital, without external venture funding, and did not disclose revenue. This approach of focusing on research free from external pressure is interpreted as a deliberate choice to avoid being bound by a short-term commercialization timetable.

The pivot to fundraising (2026). In April 2026, Liang used his own money to increase DeepSeek's registered capital by 50%, from RMB 10 million to RMB 15 million; his personal contribution rose from RMB 100,000 to RMB 5.1 million, raising

his effective control to about 84.3%. DeepSeek subsequently pursued its first external funding round of USD 3–4 billion at a valuation of roughly USD 50 billion (around USD 45 billion per some reports), led by China's national semiconductor and AI fund, with Tencent and Hillhouse among those discussing participation. Liang holds about 90% of the company, and the primary reason for accepting outside investment was to offer equity to employees in response to talent poaching by rivals.

3. The Live Test: Alpha Arena Results

In October 2025, in the real-money trading competition Alpha Arena hosted by the financial-AI lab Nof1.ai, six models (Qwen3 Max, DeepSeek Chat V3.1, GPT-5, Gemini 2.5 Pro, Claude Sonnet 4.5, Grok 4) each autonomously traded cryptocurrency perpetual futures on Hyperliquid with USD 10,000.

Correction of the figures

The winner was Qwen3 Max (about 22.32% return); DeepSeek came second at about 4.89%, while the four U.S. models recorded losses of 30.81% to 62.66%. DeepSeek reached a peak of +125% mid-competition but then gave back a large portion, ending with a single-digit final return. In other words, "an overwhelming first place" is not accurate; "a mid-race peak followed by a sharp pullback to a final second place" is correct. That said, one can observe in this the hedge-fund-like character of parent company High-Flyer, which pursues strong alpha in rising markets.

The outcome — Chinese models, trained on essentially unlimited financial-trading data, on top, and U.S. models broadly in the red — matches the result. The competition revealed that high scores on static academic benchmarks (MMLU, GPQA, etc.) do not guarantee survival and profit in highly uncertain real markets. According to the organizers' commentary, Qwen3 Max and DeepSeek managed leverage and hedging relatively stably, whereas GPT-5, Gemini, and Claude suffered heavy losses from excessive leverage and inadequate risk management.

4. The Evaluation Standard — Falsifiability, Not Accuracy

The quality of financial research is measured by falsifiability rather than accuracy. "The market may go up or down" is 100% correct but has zero information value. Genuine insight comes from concretely specifying "the reasons not to buy right now" and stating the conditions under which that logic would be wrong. DeepSeek's reasoning structure is designed to generate and re-evaluate counter-arguments to its own conclusions, so that alongside a "do not buy" conclusion it must also present a contrary scenario such as "EPS continuing to grow at double digits."

Simplifying the positioning among models: ChatGPT (the GPT family) is strong at explanation drawn from a vast knowledge base; Claude (Anthropic) is strong at advice centered on safety and alignment; Gemini (Google) is strong at information retrieval built on real-time data and ecosystem integration. DeepSeek positions "structured skepticism" as its differentiator.

5. Why Finance? — An Asymmetric Domain

Financial markets are a brutal testbed for AI. A predictive edge of just 51% can be enough to beat the market, yet it is an asymmetric structure in which a single tail risk can wipe out many successes. In this environment, the meaningful goal is not prediction but scenario calibration.

For instance, rather than delivering data showing the equity risk premium (ERP) is at dot-com-bubble levels as a mere warning, it is the ability to simultaneously evaluate "the conditions under which the market could still rise further from that level" and "the triggers that lead to collapse." Room for excess returns opens up when narrative fallacies such as "midterm-election years always rose" can be systematically dismantled.

6. Open Source and Verifiability

Black-box models are hard to use as a basis for asset allocation because their reasoning paths cannot be verified. Even while taking outside funding, DeepSeek maintains its stance of releasing open-source models and prioritizing fundamental research over short-term commercialization. When the thinking process is exposed, users can evaluate the soundness of the argument rather than just the result, which also dovetails with the procedural verification (auditability) demanded in regulated environments. The low inference cost based on MoE is a secondary benefit.

Closing

The advance of market research has come from better questions, not more accurate answers. Insight begins the moment one asks not "should I buy now?" but "is the current decline political noise or structural damage?" The grounds for DeepSeek's advantage in financial reasoning are clear.

They are: parent company High-Flyer's long experience in algorithmic trading, a history of independence that allowed it to focus on research free from outside capital pressure, and verifiability based on open source. That said, as the Alpha Arena results show, the real-world edge is not absolute (a final second place, a mid-race peak followed by a sharp pullback) and hinges on consistency in risk management. I believe this part can be overcome by revising strategy through LLM prompts.

The core lesson the market teaches is not computational speed or data volume, but that the intellectual honesty to refute oneself is what determines long-term survival.

This column is a general analysis based on DeepSeek's design architecture, publicly available information about parent company High-Flyer Quant, and public test results such as Alpha Arena. It does not recommend any specific investment.

Key References

- **High-Flyer 2025 returns & AUM:** [SCMP](#) • [Bloomberg](#) • [Hedgeweek](#)
- **Liang Wenfeng / High-Flyer founding & GPU acquisition:** [Fortune](#) • [Wikipedia – High-Flyer](#)
- **Alpha Arena final results (Qwen 1st, DeepSeek 2nd):** [The China Academy](#) • [iWeaver AI](#) • [Bitget News](#)
- **DeepSeek registered-capital increase & external funding:** [Yicai Global](#) • [TechFundingNews](#) • [The AI Insider](#)
- **Open source / AGI stance:** [TNW](#) • [Bloomberg](#)

🔗 **Related repository:** [vibe-investing](#) — an AI-driven investment-research curation combining quant theory, Python backtesting, and Claude prompt templates